

Issues Faced

Problem 1:

Cross-Origin Resource Sharing (CORS) Error

Issue:

When calling FastAPI endpoints from Streamlit (frontend → backend), browsers may block the request due to **CORS restrictions**.

Example error in browser console:

Access to fetch at 'http://localhost:8000/chat' from origin 'http://localhost:8501' has been blocked by CORS policy

Solution:

Enable CORS in FastAPI:

```
from fastapi.middleware.cors import CORSMiddleware
```

```
app.add_middleware(
    CORSMiddleware,
    allow_origins=["*"],
    allow_credentials=True,
    allow_methods=["*"],
    allow_headers=["*"],
)
```

Problem 2:

File Upload Parsing Error

Issue:

`extract_text(file)` may throw errors if the uploaded file is not properly passed to the backend.

`UploadFile` in FastAPI gives a file-like object, but `uploaded_file.getvalue()` in Streamlit gives raw bytes.

Solution:

Convert file correctly before sending:

```
files = {"file": ("document.pdf", uploaded_file.getvalue(), "application/pdf")}
response = requests.post("http://localhost:8000/summarize", files=files)
```

Problem 3:

Missing Dependency Errors

Issue:

Running the app may fail with errors like:

```
ModuleNotFoundError: No module named 'ollama_client'
ModuleNotFoundError: No module named 'doc_parser'
```

Solution:

Ensure you create `ollama_client.py` and `doc_parser.py` with proper functions (`get_chat_response`, `get_summary`, `extract_text`).

Install required libraries:

```
pip install fastapi uvicorn streamlit pydantic requests pypdf2
```

Problem 4:

Uvicorn Server Not Running

Issue:

Streamlit frontend won't work unless FastAPI backend is running at `http://localhost:8000`.

Solution:

Start backend first:

```
uvicorn main:app --reload --port 8000
```

Then run Streamlit frontend:

```
streamlit run app.py
```

Problem 5:

Large PDF Processing Crash

Issue:

Uploading large PDFs may cause memory errors or timeouts in `extract_text`.

Solution:

Use chunk-based PDF parsing (PyPDF2, pdfplumber, or `langchain.text_splitter`).
Apply summarization **incrementally per chunk**.

```
from langchain.text_splitter import RecursiveCharacterTextSplitter
splitter = RecursiveCharacterTextSplitter(chunk_size=1000, chunk_overlap=200)
chunks = splitter.split_text(content)
```

Problem 6:

Inconsistent LLM Responses

Issue:

`get_summary` or `get_chat_response` may return empty/irrelevant outputs depending on model or prompt.

Solution:

Add **prompt engineering** for clarity.

Example for summarization:

```
prompt = f"Summarize the following document in bullet points:\n\n{content}"
summary = get_summary(prompt)
```

Problem 7:

Streamlit Blocking UI

Issue:

When backend response takes time (large docs or long chats), Streamlit UI looks frozen.

Solution:

Use Streamlit's spinner and success indicators:

```
with st.spinner("Processing..."):
    response = requests.post(...)
st.success("Done!")
```